

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('bmw.csv')

df.shape
(10781, 9)

df.head(10)

      model  year  price transmission  mileage  fuelType  tax  mpg \
0    5 Series  2014   11200  Automatic    67068  Diesel  125  57.6
1    6 Series  2018   27000  Automatic   14827  Petrol  145  42.8
2    5 Series  2016   16000  Automatic   62794  Diesel  160  51.4
3    1 Series  2017   12750  Automatic   26676  Diesel  145  72.4
4    7 Series  2014   14500  Automatic   39554  Diesel  160  50.4
5    5 Series  2016   14900  Automatic   35309  Diesel  125  60.1
6    5 Series  2017   16000  Automatic   38538  Diesel  125  60.1
7    2 Series  2018   16250     Manual   10401  Petrol  145  52.3
8    4 Series  2017   14250     Manual   42668  Diesel   30  62.8
9    5 Series  2016   14250  Automatic   36099  Diesel   20  68.9

      engineSize
0            2.0
1            2.0
2            3.0
3            1.5
4            3.0
5            2.0
6            2.0
7            1.5
8            2.0
9            2.0

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10781 entries, 0 to 10780
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   model            10781 non-null   object 
 1   year             10781 non-null   int64  
 2   price            10781 non-null   int64  
 3   transmission     10781 non-null   object 
 4   mileage          10781 non-null   int64  
 5   fuelType          10781 non-null   object 
 6   tax              10781 non-null   int64 

```

```
7    mpg          10781 non-null  float64
8    engineSize   10781 non-null  float64
dtypes: float64(2), int64(4), object(3)
memory usage: 758.2+ KB

pd.isnull(df)

      model  year  price  transmission  mileage  fuelType  tax
mpg \
0     False  False  False        False  False  False  False
False
1     False  False  False        False  False  False  False
False
2     False  False  False        False  False  False  False
False
3     False  False  False        False  False  False  False
False
4     False  False  False        False  False  False  False
False
...
...
10776  False  False  False        False  False  False  False
False
10777  False  False  False        False  False  False  False
False
10778  False  False  False        False  False  False  False
False
10779  False  False  False        False  False  False  False
False
10780  False  False  False        False  False  False  False
False

      engineSize
0            False
1            False
2            False
3            False
4            False
...
...
10776        False
10777        False
10778        False
10779        False
10780        False

[10781 rows x 9 columns]

# check full null Values
pd.isnull(df).sum()
```

```

model          0
year           0
price          0
transmission   0
mileage         0
fuelType        0
tax             0
mpg            0
engineSize     0
dtype: int64

df.dropna(inplace=True)

print("\nStatistical Summary:")
print(df.describe())

Statistical Summary:
      year      price      mileage      tax
mpg \
count  10781.000000  10781.000000  10781.000000  10781.000000
10781.000000
mean   2017.078935  22733.408867  25496.986550  131.702068
56.399035
std    2.349038    11415.528189   25143.192559   61.510755
31.336958
min   1996.000000   1200.000000    1.000000    0.000000
5.500000
25%   2016.000000   14950.000000   5529.000000   135.000000
45.600000
50%   2017.000000   20462.000000   18347.000000   145.000000
53.300000
75%   2019.000000   27940.000000   38206.000000   145.000000
62.800000
max   2020.000000   123456.000000  214000.000000   580.000000
470.800000

      engineSize
count  10781.000000
mean    2.167767
std     0.552054
min    0.000000
25%    2.000000
50%    2.000000
75%    2.000000
max    6.600000

# Data Cleaning
df = df.drop_duplicates()

```


```
# Handle missing values
df = df.fillna()    # forward fill (new recommended syntax)

print("\ncleaned Dataset Info:")
print(df.info())


cleaned Dataset Info:
<class 'pandas.core.frame.DataFrame'>
Index: 10664 entries, 0 to 10780
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   model        10664 non-null   object  
 1   year         10664 non-null   int64  
 2   price        10664 non-null   int64  
 3   transmission 10664 non-null   object  
 4   mileage       10664 non-null   int64  
 5   fuelType      10664 non-null   object  
 6   tax           10664 non-null   int64  
 7   mpg           10664 non-null   float64
 8   engineSize    10664 non-null   float64
dtypes: float64(2), int64(4), object(3)
memory usage: 833.1+ KB
None

if 'price' in df.columns:
    plt.figure(figsize=(6,4))
    plt.hist(df['price'], bins=30)
    plt.title("Distribution of Car Prices")
    plt.xlabel("Price")
    plt.ylabel("Count")
    plt.show()
```


Distribution of Car Prices




```
if 'model' in df.columns:  
    plt.figure(figsize=(10,5))  
    df['model'].value_counts().plot(kind='bar')  
    plt.title("Count of BMW Models")  
    plt.xlabel("Model")  
    plt.ylabel("Count")  
    plt.xticks(rotation=45)  
    plt.show()
```




```
if 'price' in df.columns and 'model' in df.columns:
    plt.figure(figsize=(10,5))
    df.groupby('model')['price'].mean().sort_values().plot(kind='bar')
    plt.title("Average Price by BMW Model")
    plt.xlabel("Model")
    plt.ylabel("Average Price")
    plt.xticks(rotation=45)
    plt.show()
```



```
if 'mileage' in df.columns and 'price' in df.columns:  
    plt.figure(figsize=(7,5))  
    plt.scatter(df['mileage'], df['price'])  
    plt.title("Price vs Mileage")  
    plt.xlabel("Mileage")  
    plt.ylabel("Price")  
    plt.show()
```




```
if 'engineSize' in df.columns and 'price' in df.columns:  
    plt.figure(figsize=(7,5))  
    plt.scatter(df['engineSize'], df['price'])  
    plt.title("Price vs Engine Size")  
    plt.xlabel("Engine Size")  
    plt.ylabel("Price")  
    plt.show()
```



```
import seaborn as sns

if 'price' in df.columns and 'model' in df.columns:
    plt.figure(figsize=(12,6))
    sns.boxplot(x='model', y='price', data=df)
    plt.title("Price Distribution by Model")
    plt.xticks(rotation=45)
    plt.show()
```



```


import seaborn as sns
import matplotlib.pyplot as plt

# Select only the numeric columns
numeric_df = df.select_dtypes(include=['int64', 'float64'])


# Compute correlation
corr_matrix = numeric_df.corr()

# Plot heatmap
plt.figure(figsize=(8,5))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap (Numeric Columns Only)")
plt.show()

```




```
if 'year' in df.columns and 'mileage' in df.columns:
    plt.figure(figsize=(7,5))
    plt.scatter(df['year'], df['mileage'])
    plt.title("Mileage vs Year")
    plt.xlabel("Year")
    plt.ylabel("Mileage")
    plt.show()
```




```
if 'year' in df.columns and 'price' in df.columns:  
    plt.figure(figsize=(8,5))  
    df.groupby('year')[['price']].mean().plot(kind='line', marker='o')  
    plt.title("Average Price Over the Years")  
    plt.xlabel("Year")  
    plt.ylabel("Avg Price")  
    plt.grid()  
    plt.show()
```

Average Price Over the Years




```
if 'mileage' in df.columns:  
    plt.figure(figsize=(7,4))  
    plt.hist(df['mileage'], bins=30)  
    plt.title("Mileage Distribution")  
    plt.xlabel("Mileage")  
    plt.ylabel("Count")  
    plt.show()
```

Mileage Distribution



```
if 'engineSize' in df.columns:  
    plt.figure(figsize=(7,4))  
    plt.hist(df['engineSize'], bins=20)  
    plt.title("Engine Size Distribution")  
    plt.xlabel("Engine Size")  
    plt.ylabel("Count")  
    plt.show()
```



```
# Key insights Summary
def insight(text):
    print("-", text)

print("\n*KEY INSIGHTS:")
insight("Newer cars tend to have higher prices (positive correlation).")
insight("Higher mileage cars usually have lower market value.")
insight("Strong correlation found between car year, mileage, and price.")
insight("Price distribution is skewed depending on luxury model variants.")

*KEY INSIGHTS:
- Newer cars tend to have higher prices (positive correlation).
- Higher mileage cars usually have lower market value.
- Strong correlation found between car year, mileage, and price.
- Price distribution is skewed depending on luxury model variants.

df.to_csv("cleaned_bmw_dataset.csv", index=False)
print("\nCleaned dataset saved as: cleaned_bmw_dataset.csv")

Cleaned dataset saved as: cleaned_bmw_dataset.csv

import matplotlib.pyplot as plt
```

```

# Check if 'price' column exists to avoid runtime errors
if 'price' in df.columns:

    # Create figure for better readability
    plt.figure(figsize=(6,4))

    # Plot histogram to understand price distribution
    plt.hist(df['price'], bins=30)

    # Add chart title and axis labels
    plt.title("Distribution of Car Prices")
    plt.xlabel("Price")
    plt.ylabel("Count")

    # Display the plot
    plt.show()


"""

INSIGHTS:
1. This histogram shows how car prices are distributed across the dataset.
2. A higher bar height indicates more cars within that price range.
3. If most bars are concentrated on the lower price side, the data is right-skewed,
   meaning budget cars dominate the market.
4. Very high price values appearing at the extreme right indicate outliers,
   which can negatively impact linear regression.
5. This analysis helps decide whether price transformation or outlier handling
   is required before building a regression model.

"""

else:
    print("Column 'price' not found in dataset")

```



```

import matplotlib.pyplot as plt

# Check column existence for safe execution
if 'price' in df.columns:

    # Create a larger and clearer chart
    plt.figure(figsize=(8,5))

    # Plot price distribution
    plt.hist(df['price'], bins=30)

    # Titles and labels
    plt.title("Car Price Distribution")
    plt.xlabel("Car Price")
    plt.ylabel("Number of Cars")

    # Show chart
    plt.show()


    # Clean insights for project explanation
    """
    INSIGHTS:
    • The chart shows how car prices are distributed across the dataset.
    • Most cars fall within the lower to mid price range, indicating a right-skewed distribution.
    • A small number of cars appear at very high prices, suggesting
    """

```

```

the presence of outliers.
    • Such skewness and outliers can influence linear regression
results.
    • This analysis helps decide whether price transformation or
outlier treatment is needed
        before building a predictive model.
    """
else:
    print("The column 'price' is not available in the dataset.")

```



```

import matplotlib.pyplot as plt
import numpy as np

# Check column existence
if 'price' in df.columns:

    # Calculate statistics
    mean_price = df['price'].mean()
    median_price = df['price'].median()

    # Create larger, clearer figure
    plt.figure(figsize=(9,5))

    # Plot histogram

```

```

plt.hist(df['price'], bins=30)

# Mean and Median lines
plt.axvline(mean_price, linestyle='--', linewidth=2, label='Mean Price')
plt.axvline(median_price, linestyle='-', linewidth=2, label='Median Price')

# Titles and labels
plt.title("Distribution of Car Prices with Mean and Median")
plt.xlabel("Car Price")
plt.ylabel("Number of Cars")

# Legend
plt.legend()

# Show plot
plt.show()

"""
GRAPH INSIGHTS:
1. The histogram displays how car prices are distributed across the dataset.
2. Taller bars represent price ranges with a higher number of cars.
3. The concentration of bars on the lower price side indicates that most cars are budget to mid-range.
4. The long tail on the right side shows the presence of high-priced cars, confirming a right-skewed distribution.
5. The mean price lies to the right of the median, which further confirms positive skewness caused by expensive outliers.
6. These outliers can impact linear regression, making preprocessing steps such as transformation or capping important.
"""

else:
    print("Column 'price' not found in the dataset.")

```

Distribution of Car Prices with Mean and Median

